

Response to Reviews

Submission: EPDS-D-20-00186

Manuscript: Allotaxonometry and rank-turbulence divergence: A universal instrument for comparing complex systems **Authors:** Dodds et al.

Assessment from the Associate Editor:

The authors propose a very interesting method. We're sorry this review took a long time, but eventually we were able to secure three highly competent Reviewers. Reviewer 1 presents a few major reviewer comments and recommended a major revision. Upon closer inspection, however, those reviewer comments are mostly interpretational, but addressing them will make the manuscript more readily accessible to a broader readership. The reviewer comments of Reviewer 1 are in line with the reviewer comments from Reviewer 2, who also suggests the authors make a few minor clarifications. Reviewer 3 raises no criticism. Taken together, all reviews are very positive and only minor clarifications are requested.

Overall remarks concerning our response to reviewers:

1. We deeply apologize for the delay in our response. The lead author has had a series of unfortunate maladies over a full year. This paper represents an enormous effort, and our response has required sustained time to engage with fully.
2. We greatly thank the editor for their efforts and their overall positive evaluation.
3. We greatly appreciate the helpful comments from their reviewers, and for their positive feedback. Their input has led to a number of substantive improvements.
4. We have taken substantive time to revise the manuscript to properly address all concerns put forward by the reviewers, with individual responses below.
5. For ease of reading, we have framed the comments from the reviewers in light blue boxes. We also show relevant edited excerpts from the updated manuscript in light gray boxes.
6. We have expanded our description of the allotaxonographs in Figs. 1 and 2. We have added overlaid dashed boundaries with labels to their many parts (10 parts for Fig. 1, and 6 parts for Fig. 2). We have extracted these (small) labeled graphical elements and then incorporated them in the text to better facilitate our explanation. We hope that this new structure will be acceptable for EPJ Data Science.
7. We have also generally improved the manuscript with some minor edits throughout.

8. We note that our script for creating allotaxonographs also accommodates Probability Turbulence Divergence (PTD, described in a separate paper) and entropy-based divergences, e.g., Jenson-Shannon divergence (JSD) and a natural generalization.
9. We have added three more example allotaxonographs in response to reviewer's comments.
10. We have shared more code and data through the paper's Gitlab repository, and will aim to continue to build there. We intend that the paper's figures should all be rebuildable from scratch.
11. We have replaced Flipbook 3 which was a series of bare allotaxonographs for Twitter, rather than an intended series of allotaxonometric graphs.
12. We have also adjusted some terminology for clarity, and to remove unneeded attributions to individual scientists. For example, we have largely replaced "Zipf distributions" with the more plainspoken "size rankings", "size-rank distributions", and "rankings". (We still cite Zipf appropriately.)

Overall, and while we have expended considerable time and effort to respond fully and faithfully to the reviewers, we have introduced no major changes in the manuscript. As the editor indicated, we have found that only clarification has been requested.

We note that that we received no comments about the mathematical construction of rank-turbulence divergence, the core invention of the paper. As we have sought to demonstrate with our paper, we have extensively demonstrated how allotaxonographs and RTD perform for a range of disparate systems. The core instrument remains the same as do all original figures, except for Fig. 10 (was Fig. 7) which now includes more information.

We hope that our revised manuscript is suitable for publication in EPJ Data Science.

Response to Reviewer #1:

Comment 1 by Reviewer 1:

Reviewer #1: Thank you for the opportunity to review this paper. It proposes new methods to compare either multiple systems or a single system over time, including graphs and a measure of divergence. The proposed methods are demonstrated on a variety of data examples.

Major Reviewer comments:

Can we use the rank-rank histogram to detect differences in the number of component types between the two systems? For example, can we use the histogram to figure out whether there are more species in 1985 vs. 2015 in Figure 3?

As is, we do not make the total number of types explicit in the rank-rank histograms. We have intended that our instrument will work on bare rankings.

To further help make clear that rank-divergence works purely on rankings we have adjusted a paragraph from the introduction to read:

We will consider systems where each component type τ has at least one measurable—and hence rankable—“size” s_τ . We make clear that we may have no knowledge of these component sizes—we may only have rankings (e.g., book rankings provided by a seller, but with numbers of sales withheld). Our core instrument functions only on rankings, incorporating size data (if available) for minor diagnostics.

That said, and as we note below, various absolute numbers are returned within the Matlab environment when an allotaxonograph is generated (and can be computed separately).

We provide three points.

1. In general, our rank-rank histograms will strongly visually indicate if two systems are imbalanced in their respective number of distinct types (on this topic, see also our response to Question 5 from Reviewer 2).

For example, in Figs. 4 and 5, the allotaxonographs for baby girl and boy names show a distinct line separated from the main histogram only on the 2018 (right hand) side. Names that fall along this separated line are names that appear in only 2018 ('2018-exclusive types'). So without the exact numbers of types, a user (with some familiarity) will immediately see that there are more types of baby names in 2018.

By contrast, the allotaxonograph for Twitter in Fig. 1 shows a fairly balanced number of types. On both sides, we see two similarly separated lines from the main histogram, corresponding to 1-grams that appear 0 times in the comparison corpus (bottommost outer separated line) or 1 time (inner separated line).

2. The script for generating allotaxonographs (figallotaxonometer9000.m in Matlab provided in the paper's Gitlab repository) returns some diagnostics including a set of absolute numbers (this output can be ignored).

Here's an example of part of the output, using the 1968–2018 baby girl name comparison from Fig. 4, and with comments added:

```
details =
struct with fields:

Nunion: 22758    %% # of types in union of both systems
Nshared: 3552    %% # of types in intersection of both systems
N1: 8195    %% # of types in system 1 (baby girl names, 1968)
N2: 18115    %% # of types in system 2 (baby girl names, 1968)
N1exclusive: 4643    %% # of exclusive types in system 1
N2exclusive: 14563    %% # of exclusive types in system 2
totalcounts1: 1640302    %% Total # of girls in system 1 (token count)
totalcounts2: 1696917    %% Total # of girls in system 2 (token count)
```

3. (Note: Some form of the text below now appears in the manuscript. See our answer to reviewer 1's third comment below.)

The balances at the bottom right of each allotaxonograph (now labeled 1G in Fig. 1) present various *relative* comparisons derived from absolute numbers of types, and if available, 'sizes' of those types. Size is often a form of counting (words, baby names) but could be anything and may not be normalizable (e.g., market cap for corporations).

Again for baby girl names in Fig. 4, here's the extracted display of balances:

Balances:
49.2% total counts 50.8%
36.0% all names 79.6%
56.7% exclusive names 80.4%

Each allotaxonograph will display two or three bars. Per the reviewer's interest, comparisons for the numbers of types in each system are recorded in the bottom two bars, and the top one, described below, is rendered optionally.

Counts of names (and more generally, sizes of types) are not incorporated in the two bottom bars.

The second balance from the bottom shows the fraction of types in each system as a percentage of the union of types from both systems. For baby girl names, the computation is:

```
>> details.N1 / details.Nunion * 100
8195 / 22758 * 100
ans =
36.0093
```

```

>> details.N2 / details* 100.Nunion * 100
18115 / 22758 * 100
ans =
79.5984

```

These numbers correspond to the middle of the balances: 36.0% versus 79.6%.

The bottom balance shows that given each side's set of types, what percentage appears only on that side. The computation is:

```

>> details.N1exclusive / details.N1 * 100
4643 / 8195 * 100
ans =
56.6565
>> details.N2exclusive / details.N2 * 100
14563 / 18115 * 100
ans =
80.3919

```

And these numbers are shown in the bottom bar as 56.7% and 80.4%.

Lastly, a third bar will appear on top if we have measures of type size. For baby names, we have counts and the computation is:

```

>> sumtc = details.totalcounts1 + details.totalcounts2
sumtc =
3337219
>> details.totalcounts1 / sumtc * 100
1640302 / 3337219 * 100
ans =
49.1518
>> details.totalcounts2 / sumtc * 100
1696917 / 3337219 * 100
ans =
50.8482

```

The topmost of the three bars shows the relative balances for the two years: 49.2% versus 50.8%.

What the top balance bar shows is that the number of baby girls registered with social security has increased only very slightly from 1968 to 2018. The expansion of the 'lexicon' of baby names is thus a strong one and not an artifact of increases in sheer numbers. We discuss this further in the manuscript.

We have added the following paragraph to the baby names section:

We emphasize that the balance indicators are for baby names appearing at least five times. For our present work, and in attempting to maintain

uniformity across allotaxonographs, we do not attempt to adjust for names appearing less than 5 times, though this would be possible for the topmost balance for total counts given we have that information separately. Clearly the balance values would shift if we had complete data sets for baby names; estimating errors for these estimates would be meaningful future work.

Comment 2 by Reviewer 1:

Also, can we use the graph to gauge the difference in size between the component type of rank r in System 1 and the component type of System 2 of the same rank?

As per our response to the question above, for RTD, we wanted to build a tool that only works with ranked data. While in practice, we will have systems where type ‘sizes’ are known, the intent here was to build a fully ranked-based instrument in the spirit of well established rank-based statistical methods (e.g., the Spearman coefficient).

We add two points:

1. The main allotaxonometer script can produce allotaxonographs that compare by only type size. These outputs will show only an allotaxonograph with an annotated histogram and no overlaying measure of difference.
2. In our second paper on allotaxonomety, we introduce and explore probability-turbulence divergence (PTD), where we make size (as measured by probability or rate) explicit. We note that the allotaxonometer code does include the ability to compare systems where we have sizes that are not counts or probabilities or cannot be interpreted as relative fractions in some way (e.g., heights of buildings). We construct PTD with a tunable parameter $0 \leq \alpha < \infty$ in a similar fashion to RTD. We show that PTD has different limiting behaviors from RTD, and that it matches up with a collection of existence difference measures including $L^{(p)}$ norms, the Sørensen-Dice coefficient (the F_1 statistic), and the Hellinger distance.

We add that we prefer PTD over the commonly used Jenson-Shannon divergence (JSD), which we have a separate article about, not yet on the arXiv. JSD creates a mixture of two systems which does not exist (half Moby Dick, half Pride and Prejudice for example). PTD works as a true difference and is what one would want in constructing a form of “type calculus” (“lexical calculus” for text corpora). Our allotaxonograph script nevertheless will produce JSD comparison plots, along with a particular generalized entropy divergence.

The PTD manuscript can be found here: <https://arxiv.org/abs/2008.13078>

We note that we had mentioned type calculus and lexical calculus in the conclusion:

Allotaxonomety may be viewed as part of a larger analytic framework of “lexical calculus” and, more generally, “type calculus.”

We have now added “type calculus” to the abstract.

Comment 3 by Reviewer 1:

How do we interpret the percentages in the “Balances” section at the bottom of Figure 1? Are these percentages unrelated to your alpha tuning parameter?

For an overview of the Balance percentages, we point back to our response to Reviewer 1’s question 1 above.

We add here that the balances are independent of alpha, and indicate relative percentages of type sizes (e.g., word counts), types, and exclusive types.

As we mention in the overview for our reply, we have outlined 10 distinct cutouts of the allotaxonograph in Fig. 1 to help with our discussion. We have taken these cutouts and added them to the text to, we hope, greatly help the reader with being able to comprehend the instrument.

We describe balances in Sec IIB: “Rank-Rank Histograms for Basic Allotaxonomy”. The cutout for Fig. 1 for balances is Fig. 1G and can be found looking through the subsections of Sec IIB.

We reproduce the relevant text here:

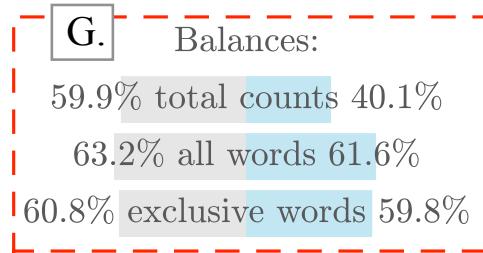


Fig. 1G: For all allotaxonographs, we show balances at the bottom right of the rank-rank histogram. Two kinds of exclusive type comparisons for the numbers of types in each system are recorded in the bottom two bars, while the top bar conveys information about type sizes. The top bar is rendered only when sizes are known (this is the only part of the allotaxonographs that is not determined purely from type rankings). In the present paper, we work from datasets with known type sizes, and all allotaxonographs will have all three bars. All balances show normalized quantities rather than absolute numbers

As we will see, these three balances can vary greatly across system comparisons.

The top balance bar shows the relative balance of the two systems’ sizes,

if known. For our Twitter example, this top bar shows the breakdown of total counts of 1-grams (type size) between the two dates at 59.9% and 40.1%. We thus see that the election generated considerably more tweets than Charlottesville.

The middle balance bar shows the fraction of types in each system as a percentage of the union of types from both systems. For Twitter, we have that of all words in the combined lexicon for the two days combined, just over 60% appear on each of the two days (63.2% and 61.6%).

The bottom balance bar shows that given a system's set of types, what percentage of those types appear only in that system. For the Twitter example, we take the separate lexicons for each day, and find that around 60% of words are exclusive for both days, further giving a sense of strong turnover (60.8% and 59.8%).

We add that the script for generating allotaxonographs (figallotaxonometer9000.m in Matlab provided in the paper's Gitlab repository) returns some diagnostics including the underlying numbers used to compute the above balances.

Comment 4 by Reviewer 1:

The paper states that the standard methods lack transparency and flexibility. It may be helpful to apply some standard methods in one of your case studies to demonstrate the differences from your method.

In preparing this manuscript, we had separated out work where we compare methods (e.g., JSD vs RTD vs PTD), but we had left residual guidance in the abstract. This work is still being written up (as mentioned above).

We have removed two sentences from the abstract that misleadingly pointed in the direction of other measures being deficient. We have also removed any similar statements in the main text.

Comment 5 by Reviewer 1:

How sensitive is your method to small perturbations? For example, if you fixed the errors in the Berkshire Hathaway and DowDuPont market caps in the last case study, how much do the divergence and histogram change?

In general, using ranks will give a more robust instrument, and is a key part of our overall motivation.

From the paper:

Fourth, rank orderings potentially allow for powerful and robust non-parametric statistical measures such as Spearman's rank correlation coefficient. All told, while in moving from sizes to rankings we may trade information for some simplification, we still preserve a great deal of meaningful structure.

We understand the reviewer's point, however, that we could perform a deep analysis of sensitivity. In this foundational paper, we have explored the effects of truncation only, and reserve further work for later (by us, and, hopefully, others as well). We believed (and continue to believe) that the few small clear errors in the market cap database were worth leaving in to show how outlier kinds of errors are easily observed with allotaxonographs.

The above said, we corrected the Berkshire Hathaway and DowDuPont data and recomputed $D_{1/3}^R$, finding it be unchanged to three significant figures at 0.411. The adjusted allotaxonograph (which we have added to the paper, and is included in our response below) now make Berkshire Hathaway clearly ascendent on the right side, as it was ranked 5th overall by market cap at the end of 2018.

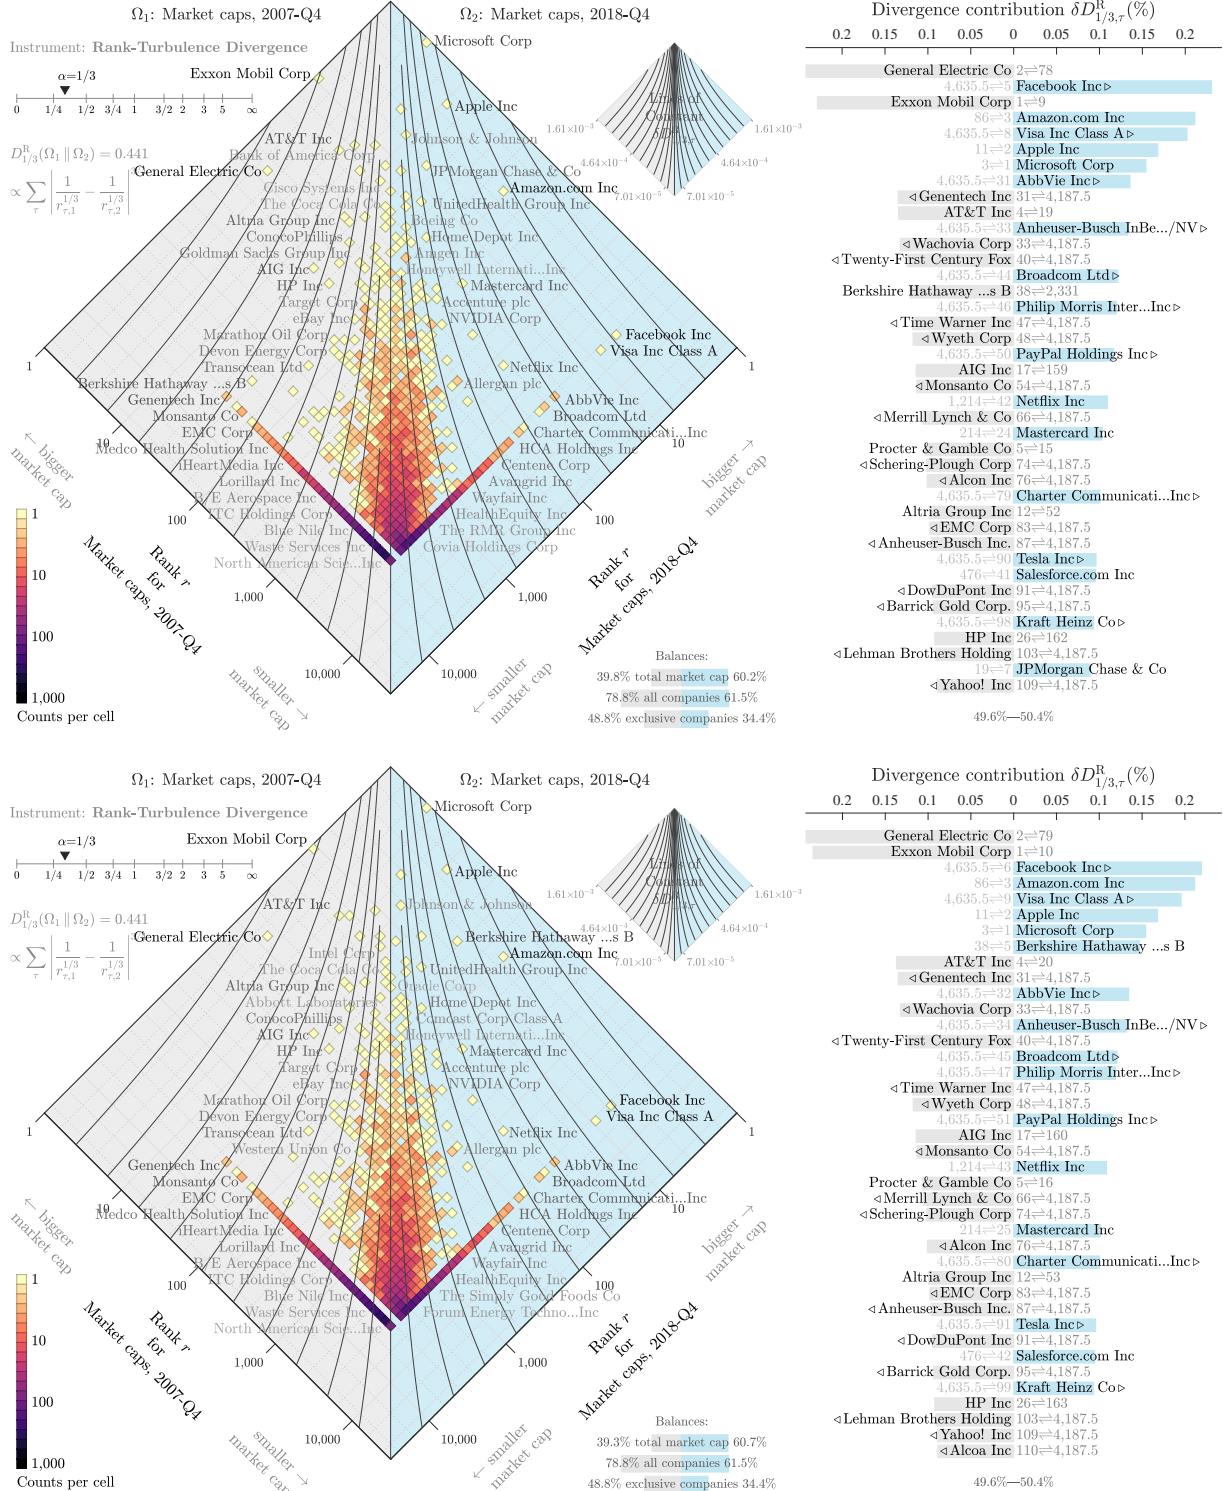
Per the above, we have also added the following to the manuscript:

The dataset for market caps does have some missing and erroneous data. DowDuPont's market cap for the last quarter of 2018 is absent and is consequently shown to have plummeted from a rank of $r=91$ in 2007 to equal-to-last in 2018. Berkshire Hathaway's market cap is clearly misrecorded for the last three quarters of the dataset (apparently dropping from \$528.33B to \$0.34B at the end of 2018).

We take the opportunity to perform a small test of the sensitivity of rank-turbulence divergence by correcting the data for these two companies. For DowDuPont, with further sourcing, we find the year-end 2007 and 2018 market caps were reported as \$37.06B and \$121.34B, and for Berkshire Hathaway, \$149.56B and \$502.37B. Upon making these corrections, we first find again that $D_{1/3}^R=0.411$, unchanged to three decimal places. both. In the corrected allotaxonograph (Fig. 9), Berkshire Hathaway's location shifts to the right side of the histogram ($r=38\rightarrow 5$) and is now listed as the 7th overall strongest contribution for $D_{1/3}^R$. DowDuPont no longer makes the top 40 of the list of contributions. While these two changes are dramatic, the remainder of the allotaxonograph remains essentially identical.

Nevertheless, we have chosen to leave such errors in Fig. 8 to help demonstrate the importance of using a rich, graphical allotaxonometric instrument. With a naive measurement of divergence, we would easily miss problematic data points. Evidently, and beyond our present paper's interests, for any further investigations, these two errors suggest that considerable effort should be made to further clean the market cap dataset

For ease of comparison, we include both the uncorrected and corrected market cap allotaxonographs below (Figs. 8 and 9) and the corrected one below.



Caption for second figure:

Allotaxonometric comparison of publicly traded US companies in 2007 and 2018 by fourth quarter market capitalization with corrections for Berkshire Hathaway and DowDuPont Inc. To be compared with Fig. 8. The 2018 market caps for both these companies were recorded incorrectly in the original data set. The revised allotaxonograph shows Berkshire Hathaway—which in fact rose to $r=5$ in 2018—now appears prominently on the right hand side of the histogram and the contribution list. DowDuPont Inc’s corrected rank for 2018 means that it did not contribute as strongly, and now no longer appears in the top 40 contribution list. Overall, the allotaxonograph is largely identical to that of Fig. ???. The rank-divergence is unchanged at $D_{1/3}^R = 0.411$.

Comment 6 by Reviewer 1:

What is the computing time to calculate the divergence and generate the histograms in your four case studies?

Generally, the calculations and graphics require a few seconds of compute time. In our experience, the only speed limitation arises when loading massive datasets (e.g., Twitter). Separate from our analyses, computational preparation of such large datasets is where we see delays. For the Twitter data, for example, enumerating n -grams for each day (20GB of compressed text) can take many CPUs running continuously to keep up with real time, with more than one million previously unseen types encountered on a given day.

The computational bottleneck in producing the RTD calculations is confined entirely to the counting of elements to produce ranks.

Comment 7 by Reviewer 1:

Minor Reviewer 1 comments:

The terms “Zipf” and “1-gram” should be defined.

We have made the initial reference to Zipf more clear, connecting only to “Zipf’s law”. We have also now defined Zipf distribution and n -grams in the manuscript. We have replaced “Zipf ranking” with “size ranking” throughout, as we do not need to name this simple kind of ranking after an individual (we of course cite Zipf and others where appropriate).

Added to the manuscript:

Here, we will refer to such “a component ranking by decreasing size” as a “size ranking” or, for brevity and given our paper’s framing, simply a “ranking”.

In the main text, we have added this definition/clarification of n -grams when the term first appears:

For Twitter, subsampling n -grams—phrases containing n consecutive words and/or other text elements—allows for robust estimation of the rates of common n -grams but not rare ones.

And in the caption of the Fig. 1:

Words are extracted first as 1-grams (contiguous sets of non-whitespace characters) from tweets identified as English . . .

Comment 8 by Reviewer 1:

In Paragraph 2 of Part B in the Introduction, it would be helpful to get a concrete example of a component type.

We have added an explanatory paragraph at the start of section B in the introduction, as well as pointed to this when we first mention size.

In “Sec. I. Introduction A. Instruments that capture complexity”, we now have:

But for complex phenomena made up of a great many types of components of greatly varying ‘size’ (clarified below in Sec. I B)—language, ecologies, stock markets—we must confront two major problems with our dashboards of simple instruments [2].

We have put size in quotes to indicate it should be construed more generally, or at least, that the reader will need to see our description.

Then in “Sec. B. Size rankings, Zipfs law, and rank turbulence” (we have removed “Zipf distributions” from the title of this section), the section begins with:

While the instrument we develop here will have broader application, its construction focuses on two regular features of complex systems: Heavy-

tailed ‘size-rank distributions’ (rather than laws), and what we will call ‘rank turbulence’—a phenomenon of system-system comparison. We describe and discuss these two common signatures of complex systems in turn.

We will consider systems where each component type τ has at least one measurable—and hence rankable—‘size’ s_τ . To help make clear what we mean by component types and component sizes, we list some examples in the realms of language, ecology, and stock markets.

Example types: ‘the’, ‘love’, and ‘spork’.

Type size: Number of times a word appears in a book.

In “Pride and Prejudice”, for example, the corresponding counts are 4,058, 90, and 0. In linguistics, the component type and component size distinction is referred to as types and tokens [10], while more abstractly in semiotics, we have the signifier and the signified.

Example type: The species *Ornithorhynchus anatinus*, the platypus.

Type size: The number of platypuses (‘instances’ of the species) living in Australia in the wild.

Given an ecology, species may be ranked by their population numbers.

Example types: The publicly traded companies of Apple and Microsoft.

Type size: Market capitalization.

Apple and Microsoft may be viewed as components of the publicly-owned corporate world. The sizes of corporations may be broken down into many rankable dimensions such as annual revenue or number of employees worldwide. Market capitalization represents a kind of current collective belief in terms of money.

The three examples above show some of the range of what size can mean, and why we cannot readily lift terminology out from one domain to apply to all. Size for a word in a corpus means the number of indistinguishable instances of that word (many identical entities—tokens); size for species means the number of ‘biological replications’ of an individual type (many genetically similar entities of varying ages); and size for a corporation means a monetary value (one entity).

Some further examples of component size are rate (e.g., appearance of words in streaming text), physical dimensions (e.g., typical animal length or weight

for a species), social popularity (e.g., number of plays of a song) scoring in sports by individual players, and so on.

We make clear that we may have no knowledge of the underlying component sizes—we may only have rankings (e.g., book rankings provided by a seller, but with numbers of sales withheld). Our core instrument functions only on rankings, incorporating size data (if known) for minor diagnostics.

Comment 9 by Reviewer 1:

Define the notation $r_{\tau,1}$ and $r_{\tau,2}$, which seems to be first used in Part B of Section II. I believe they are the ranks of component τ in Systems 1 and 2, but this should be clarified.

The reviewer reads the notation correctly, even though, as the reviewer notes, we did not formally provide a definition.

To correct this oversight, we have added the following sentence as a standalone paragraph (now the third paragraph of Sec. II B):

In general, we will denote the rank of type τ in system $\Omega^{(1)}$ by $r_{\tau,1}$, and its rank in system $\Omega^{(2)}$ by $r_{\tau,2}$.

Comment 10 by Reviewer 1:

In Figure 1, the caption has substantial overlap with the main text. Removing some of the overlap would help highlight the new information.

It is our firm preference to make the captions reasonably complete in themselves. Readers may choose to work through the paper in different ways.

Comment 11 by Reviewer 1:

In the first two paragraphs on page 6, some of the words that are described (such as “Harambe” and “voted”) do not seem to be in Figure 1.

Many thanks for pointing this out—we appreciate the reviewer’s careful reading.

Generally speaking for rarer types, each run of the allotaxonometer script may have different annotations automatically rendered. There are a number of reasons, mainly if

annotation spacing is changed, or if a particular rank-rank pair is held by more than one type. In the latter case, the annotated type will be randomly chosen, and for some comparisons this means the rendered allotaxonometer will always be different. For example, for two days on Twitter in our data set, the number of 1-grams that appear once on one day and not on the other run into the millions.

We have adjusted the language describing Fig. 1 to match words found at the edges of the histogram. ‘Gorilla’ appears in the figure, so we have been able to retain a description of Harambe for context in the main text.

Comment 12 by Reviewer 1:

Could you give some intuition of why you chose $\alpha = 0$ for the BCI example?

We make the choice of $\alpha = 0$ by inspection. In the appropriate supplementary PDF flipbook, the effect of tuning alpha can be observed. Future work, which we discuss at the end of the paper, would seek to find a method to determine an optimal alpha for specific comparisons.

A secondary aspect of the choice to use $\alpha = 0$ is simply to demonstrate how such an allotaxonograph looks. The BCI example, with its limited rank turbulence and vertical shaping within the histogram, provides a good instance where $\alpha = 0$ is at least visually reasonable. A choice of $\alpha = 0$ would make for a poor fit for the Twitter example (and for written texts in general).

We have revised and expanded the section reference by Reviewer 1.

We numerically compare the 1985 and 2015 distributions by applying rank-turbulence divergence with $\alpha = 0$, finding $D_0^R(R_1 \parallel R_2) = 0.077$. By inspection, we choose $\alpha = 0$ here because of the match of the histogram with the verticality of the contour lines (we address optimal selection of α in our concluding remarks). The nature of the BCI example affords us an opportunity to demonstrate the limit of $\alpha = 0$ for allotaxonometry, and is a secondary reason for including an example from ecology.

In Flipbook S4, we show how the allotaxonometer performs with α varying away from 0 to ∞ . The visual match on the contour lines continuously degrades.

Comment 13 by Reviewer 1:

Also, what is the $D_{0;\text{rand}}^R$ term mentioned in that example?

We have now explained what we mean by a randomized version, and explicitly connected

this explanation to the notation.

The section (which directly continues on from the excerpt above) now reads:

The BCI example's histogram is far from what we would expect of randomized systems (Fig. II). To see how RTD quantifies the difference between an observed system and the randomized version, we construct a set of pairs of randomized systems, measuring rank-turbulence divergence for each. We do this by randomly permuting the species names within each system while leaving species counts fixed, thereby keeping the size-rank distributions the same. We can perform such a randomization for any system-system comparison (and we do so below again for baby names).

We denote the average randomized divergence for two rankings R_1 and R_2 as:

$$D_{\alpha; \text{rand}}^R(R_1 \| R_2). \quad (1)$$

For the BCI example, we find that the score $D_0^R(R_1 \| R_2) = 0.077$ is well short of the randomized equivalent of $D_{0; \text{rand}}^R(R_1 \| R_2) = 0.376$ (average of 100 randomizations; standard deviation $\sigma=0.012$).

The quantity $D_{0; \text{rand}}^R(R_1 \| R_2)$ appears explicitly in the paper twice, once in the BCI section (Sec. II C), and once in the baby girl name section (Sec II D).

We have also made the first appearance of our general notation for RTD a standalone equation (now Eq. 1).

$$D_\alpha^R(\Omega_1 \| \Omega_2). \quad (1)$$

Response to Reviewer #2:

Comment 1 by Reviewer 2:

It was my pleasure to review this article introducing a new framework to compare heavy-tailed ranked lists. Specifically, the work motivates and derives a novel divergence measure for ranked lists, and then illustrates its usefulness across four examples including: Twitter word frequency, baby name popularity, species abundance, and firm sizes.

The presentation is rich and clear, elucidating the central framework and offering wonderful insights into the motivating examples and applications. The problem this article addresses is relevant to a broad range of data science applications.

Therefore, I strongly recommend to publish the article in EPJ Data Science.

We are delighted by Reviewer 2's very positive feedback.

Comment 2 by Reviewer 2:

Reviewer comments: Even though the presentation is quite clear, it may benefit from an organization that extracts a few non-technical discussions from the rest of the text - i.e. a clearly identified general audience guide to allotaxonomy and the application of rank-turbulence. This could support usage in fields outside of physics / data sciences.

Thank you.

We will author a blog post and a Twitter thread to accompany formal publication of the manuscript that will advertise the utility of the instrument to other disciplines.

We are also working on a Python and hopefully Julia version of the allotaxonometer code (the original code was developed in Matlab) to enable and encourage broader adoption.

We and others have also used allotaxonographs in a number of other papers, and we now cite some of these in the manuscript's conclusion.

Our intent is that our foundational paper be thorough and detailed, and will serve as a reference text, and we hope that the examples are illuminating and inspirational.

Comment 3 by Reviewer 2:

Relatedly, it would be helpful if the authors could elaborate on section II E and the processes used to select different values of alpha for the different applications. I'm imagining the framework being used by a non-expert who just keeps whatever default value of alpha is encoded instead of exploring the full range - or on the flip side, being overwhelmed with the many potential choices of alpha and not understanding what

features of the histograms are being accentuated by the divergence in different limits. Can you distill a few guiding principles or questions to help motivate an appropriate selection?

We have improved our discussion of the choice of α in the conclusion, as well as given this discussion its own subsection (see excerpt below).

In this initial work, we justify the choice of alpha by looking at the shape of rank-rank histogram as a guide. The rank-rank histogram also immediately shows us when type (or lexical) turbulence is not present (e.g., when the type rankings for the two systems are random shufflings of each other).

In short, our advice for readers is to use the visual help of the allotaxonographs. In some of the supplementary PDF flipbooks, the effect of tuning alpha can be observed, and readers can see how RTD behaves when the fit is good or bad.

In follow up work that has proved difficult, and which we discuss at the end of the paper, we hope to identify a computational method to determine an optimal alpha for specific comparisons.

VI C. Determination of α

In our initial work, we have made the choice of the tuning parameter for rank-turbulence divergence, α , a visually guided one. The user gains much from inspecting the rank-rank histogram alone, and, in our experience, is then readily able to choose an α for which the allotaxonometric contour lines best match the form of the histogram. A visually guided choice will be sufficient in cases of comparing two or a small number of systems.

When rank turbulence presents as a scaling law—which is regularly the case for text corpora (e.g., Twitter, books)—we would want to be able to determine an optimal α . While for generalized entropy approaches for single systems, the limit of linear scaling and Shannon’s entropy demarcate the boundary between accentuating the common or the rare [7, 33, 53, 54], we have found that for system comparisons, the optimal value of α , if it exists, is dependent on the pair of systems being compared—there is no universal value.

Generally speaking for text corpora, as two texts become more distinct, the optimal α increases. This is particularly clear for a corpus like Twitter, for example, where rank turbulence increases with the calendrical separation of two dates being studied.

We have left open the possibility of an analytic connection between the rank-turbulence scaling described at the end of Sec. I B, and, to the extent

that well-defined scaling is present, with an optimal α for rank-turbulence divergence.

Even with an optimization method for determining α , we urge readers to always look at the visuals provided by our allotaxonographs—the maps—for confirmation of fit.

Comment 4 by Reviewer 2:

The examples illustrating the effects of partial sampling were helpful, but it would be nice to see how the divergence measure itself changes with subsample size: does it converge or is it sensitive to the influx of different cross-ranks uncovered by larger subsample sizes?

First, for each of the five main case studies, we do in fact show the full allotaxonographs in the series of Flipbooks S10, S11, S12, S13, and S14 in the supplementary material.

From Sec IV. Guide to Flipbooks, here is our description of these flipbooks:

Flipbook S10—Word use on Twitter, truncated: Full series of allotaxonographs corresponding to histograms of row 1 in Fig. 7 with $\alpha = 1/3$.

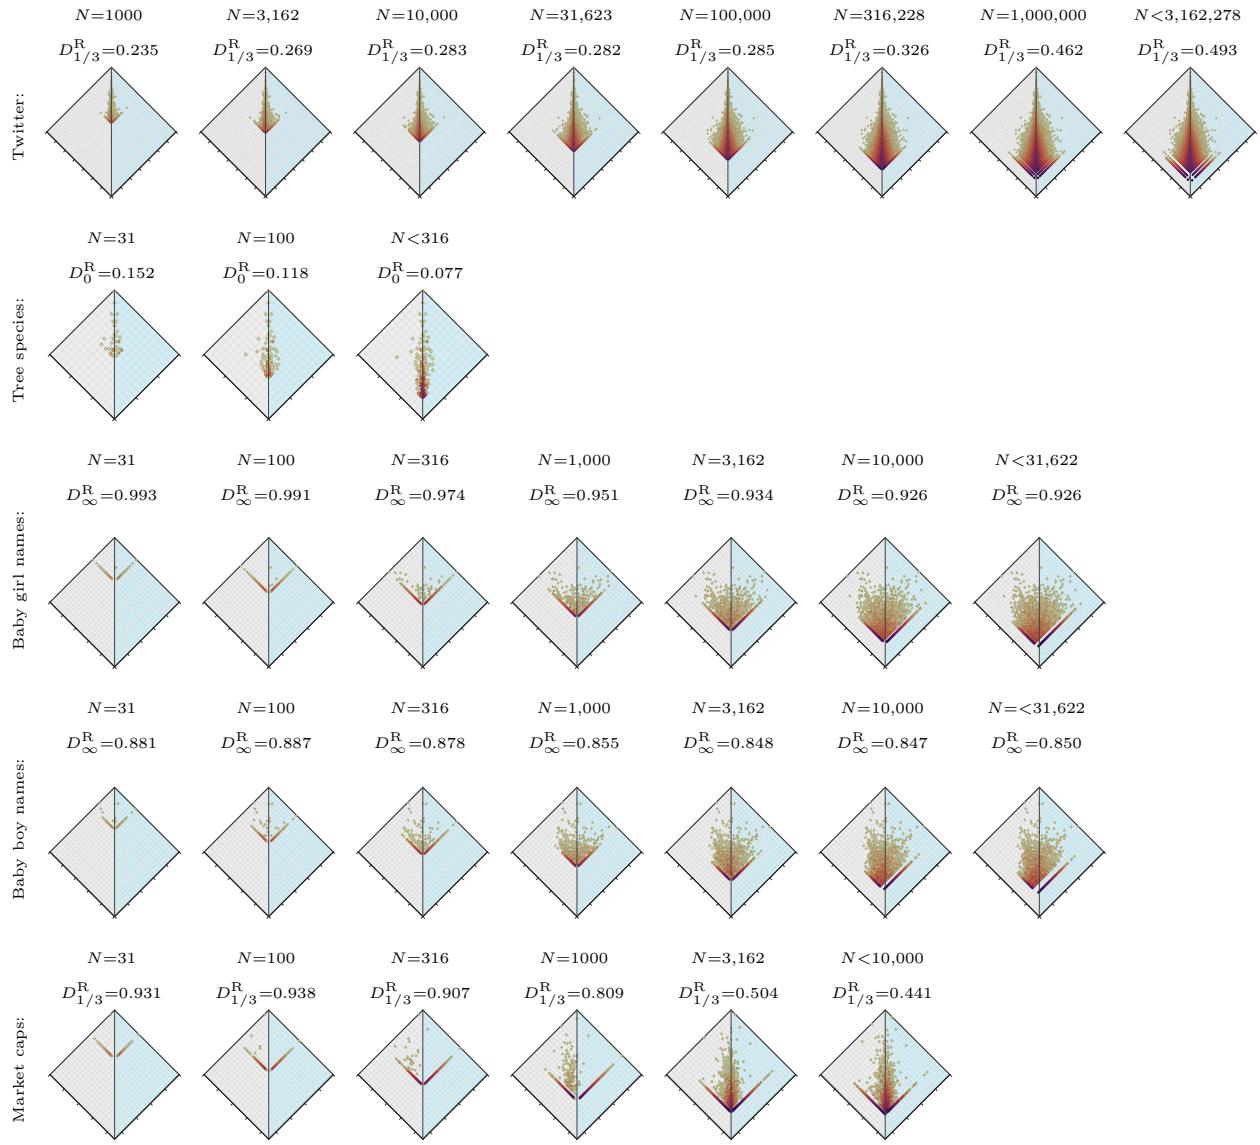
Flipbook S11—Tree species abundance, truncated: Full series of allotaxonographs corresponding to histograms of row 2 in Fig. 7 with $\alpha = 0$.

Flipbook S12—Baby girl names, truncated: Full series of allotaxonographs corresponding to histograms of row 3 in Fig. 7 with $\alpha = \infty$.

Flipbook S13—Baby boy names, truncated: Full series of allotaxonographs corresponding to histograms of row 4 in Fig. 7 with $\alpha = \infty$.

Flipbook S14—Market caps, truncated: Full series of allotaxonographs corresponding to histograms of row 5 in Fig. 7 with $\alpha = 1/3$.

Second, and however, we concur with Reviewer 2 that the figure of thumbnail allotaxonographs in the paper (now Fig. 10, was Fig. 7) could be improved by adding the rank-turbulence divergence scores, which we have now done (reproduced below). We have slightly restructured the figure to make the log spacing even for each sequence of subsamples.



We have added the following to the caption for Fig. 10:

The paths of convergence towards the divergence score vary and may be uneven if usually monotonic, depending on the systems being compared and the choice of α .

We have also restructured the caption slightly to point the reader to the flipbooks in the third sentence rather than the last.

And as the last part of our response to Reviewer 2's comment, we have added the following to the main text at the end of Section III F. "Truncation Effects for Rank-Based Allotaxonographs":

We see that the values of D_α^R for the truncation sequences approach the ‘true’ value in largely monotonic, if different, ways. The value of $D_{1/3}^R$ for the Twitter study approaches from below, deceptively exhibiting a flat section up to $N=10^6$. The ecology example starts above and moves down towards the overall score of $D_0^R=0.077$. Baby names and markets caps similarly both start above their respective overall scores for D_∞^R and $D_{1/3}^R$, and move downwards, though their scores for strong truncation are close to one as they appear as disjoint systems (their histograms are predominantly ‘vee’-shaped). While the baby name scores drop slowly and not far (0.993 and 0.881 for 31 names for girls and boys down to 0.926 and 0.850 for all names), the market cap study only starts to gain more than the ‘vee’ shape when N is into the thousands. Because the market cap data comparison is a blend of large-scale turnovers around a relatively stable core, the drop is slow and then fast and further (0.931 for 31 corporations down to 0.441 for all corporations).

Our work aside, we expect any divergence measure will likely vary as orders of magnitude more data is included. And we add that in certain circumstances, choosing to truncate a data set may be a well justified treatment of data.

Finally, we note that while some form of truncation is a common measurement issue with real data for complex systems with many components, it is certainly not the only one. Exploring how other kinds of measurement errors affect rank-turbulence divergence would be a natural area of future work.

Comment 5 by Reviewer 2:

Further, how does the divergence measure work in the presence of unbalanced list sizes, say when comparing the word usage of English tweets between countries? I noticed that some imbalance likely appears in the baby names example. Does the skewness of the underlying histogram bias the divergence and resulting understanding of turbulent names?

We have constructed rank-turbulence divergence (and subsequently probability-turbulence divergence) so that they inherently accommodate uneven systems sizes. We don’t see the measure as being a bias but rather simply capable of revealing whatever differences there are between the systems. With familiarity, the shape of the histogram along with the balance bars will give the user a rich understanding of how the systems differ (regardless of what kind of divergence is being used).

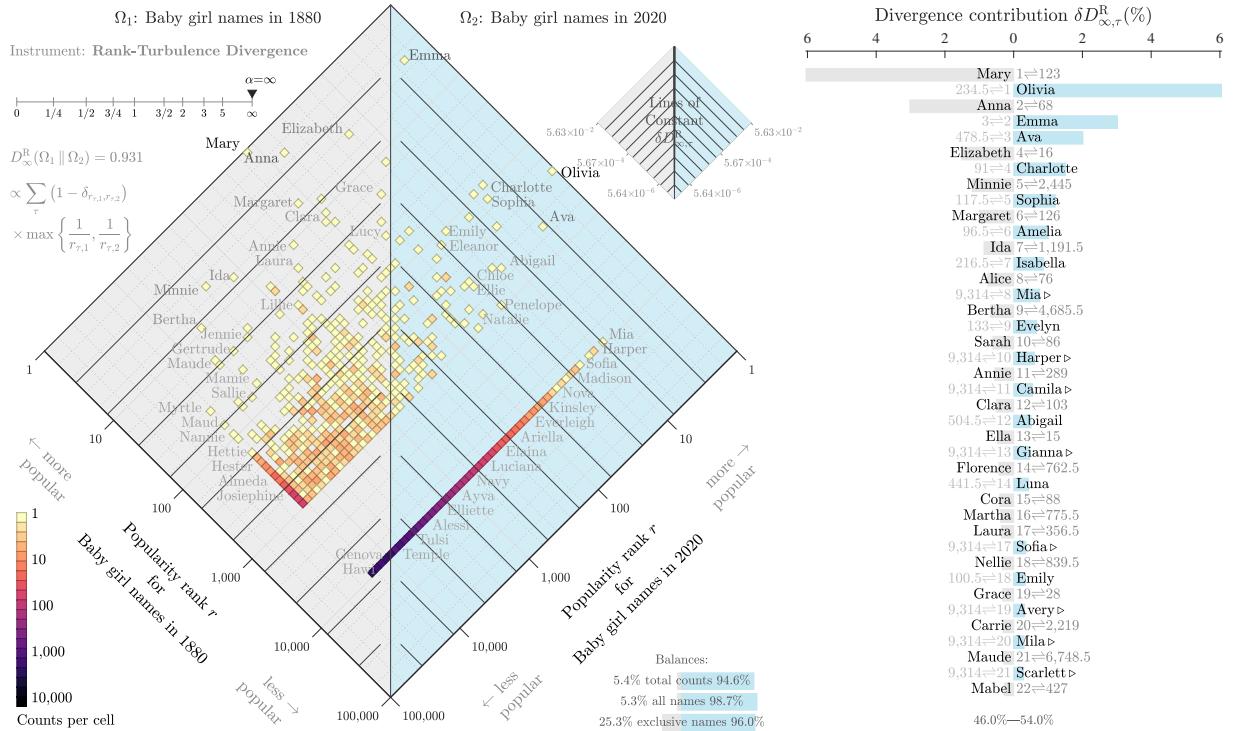
We have added two more allotaxonographs for systems of strongly different sizes as well as an accompanying discussion. We chose to extend the baby name section because baby names in 1880 versus 2020 provide a good example of size disparity between two systems.

As we note earlier in response to Reviewer 1, that the first two baby name

allotaxonographs are for roughly equal populations (1968 and 2018), and that the asymmetry of the histogram is due to the larger set of distinct names used in 2018 relative to 1968. The relevant excerpt (from the original manuscript):

We note that the asymmetries of both histograms—their apparent right-side ‘heaviness’—are not due even in part to changes in overall numbers. Using total birth numbers (see Sec. V), the total number of girl names recorded in 1968 and 2018 are comparable at 1,709,551 and 1,846,101 (7.99% increase); for boys, these numbers are 1,775,997 and 1,928,871 (8.61% increase). The number of year-exclusive names in the 1968 and 2018 are strikingly different however: 8,194 and 18,029 for girls (120% increase), and 4,742 and 14,004 for boys (195% increase). Two of the likely major factors which have lead to this explosion in name-space are immigration and a cultural shift towards parents creating novel names.

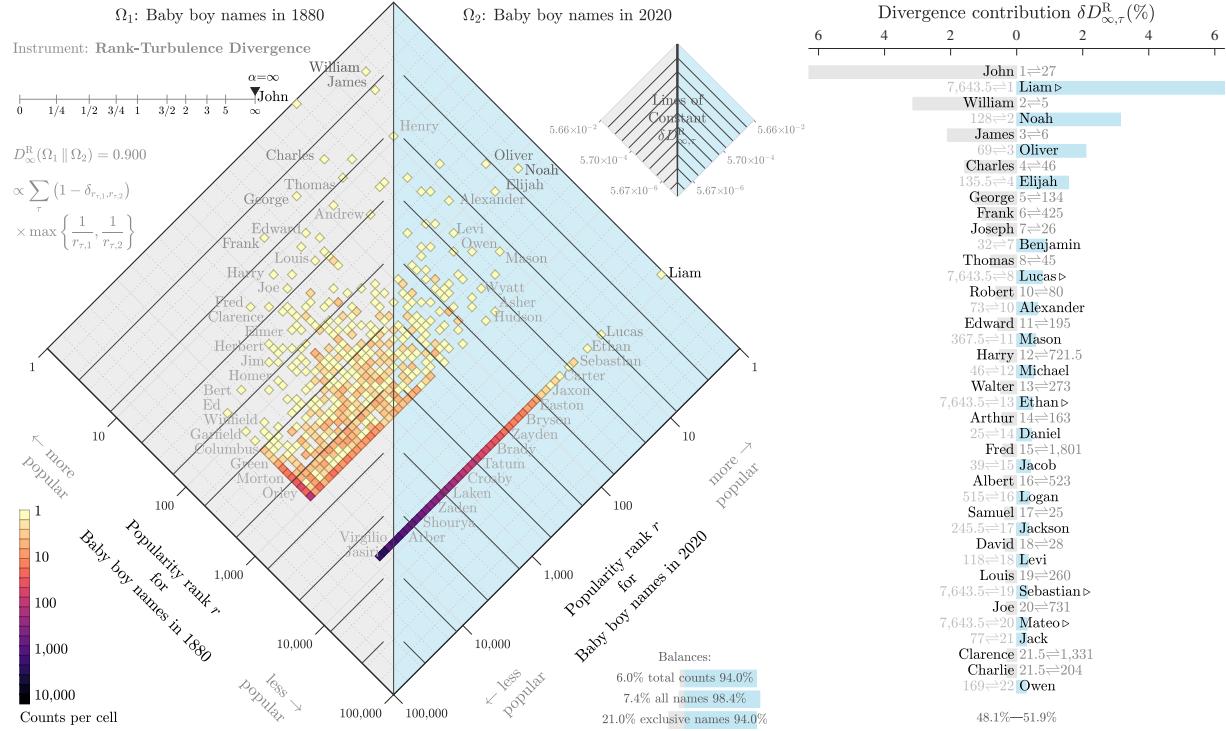
The new figures (Figs. 6 and 7) and their captions, along with the discussion excerpted follow:



Caption:

Allotaxonograph comparing US baby girl names for the years 1880 and 2020. This figure is in part a demonstration of how allotaxonographs competently perform when the sizes of two systems differ strongly. For 1968

and 2018 in Fig. 5, the balance of total baby girl names is an almost even 49.2% and 50.8%. By contrast, for 1880 and 2020, these percentages are 5.4% versus 94.6%. The choice of $\alpha = \infty$ again means that the top names for each year will dominate D_∞^R , regardless of their rank in the comparison year (unless a name has equal rank in both years).



Caption:

Allotaxonograph comparing US baby boy names for the years 1880 and 2020, companion to Fig 6. The year 1880 has 6.0% of the total baby boy names of both years combined, while 2020 has 94.0%. For 1968 and 2018 in Fig. 6, the equivalent numbers are 49.0% and 51.0%.

Added discussion.

To close out our study of baby names, we add two more allotaxonographs whose primary purpose is to show how our instrument performs when system sizes differ strongly. In Figs. 6 and 7, we compare US baby girl and boy names in 1880 and 2020, a 140 year gap.

We make some observations about balances, the rank-turbulence divergence scores, the rank-rank histograms, and the changes in naming from 1880 to

2020.

For the preceding allotaxonographs (Figs. 1–5), the largest difference for system sizes has been for Twitter in Fig. 1. The date 2016/11/09 carried 59.9% of all tweets from the two dates combined, with the other 41.1% on 2017/08/18 (top balance bar, bottom right of the histogram).

By contrast, of the total number of baby girls born in 1880 and 2020, the years separately account for 5.4% and 94.6% respectively, a factor of roughly 17-fold (new Fig. 6) For boys, these weights are similar at 6.0% and 94.0%, around 16-fold (new Fig. 7).

Because of the large increase in registered babies being born, the two kinds of type balances are consequently more extreme. For example, of the combined types for the distinct baby girl names in 1880 and 2020, only 5.3% were used in 1880, while 98.7% were used in 2020. For exclusive types, 25.3% of 1880’s distinct baby girl names appeared only in 1880, while 96.0% of 2020’s were not used in 1880.

For baby girl names, the value of $D_{\infty}^R = 9.31$ for this 140 year comparison is slightly higher than that for the 50 year gap between 1968 and 2018, $D_{\infty}^R = 9.26$, while for boys the increase is from $D_{\infty}^R = 0.850$ to 0.900.

In general, the rank-rank histograms of these disparately sized systems will show a strongly separated, highly dense line corresponding to exclusive types on the side of the large system. For both baby girl and boy names in 1880 and 2020, the separated line is around an order of magnitude from the main body of the histogram, and the component cells are high count ones. While this separation could occur for equal-sized systems if the type counts differ enough, the count density of the separated line will not be as strong. With familiarity, a glance at the balance bars will clarify these details.

Rank-turbulence divergence with $\alpha=\infty$ is a function only of the highest rank for each type (Eq. 11). As such, the main contributions for girls come from ‘Mary’ ($r=1$ to 123) ‘Olivia’ ($r=234.5$ to 1), while for boys the leaders are ‘John’ ($r=1$ to 27) and ‘Liam’ ($r=7,643.5$ to 1, not used in 1880).

For evolving complex systems, allotaxonographs can help lead us to examine time series for individual types that occupy interesting locations in the rank-rank histogram. For baby girl names, ‘Emma’ stands out as a name that was enormously popular in both 1880 ($r=3$) and 2020 ($r=2$). But the story for Emma proves to be akin to that of Vonnegut’s man-in-a-hole’s emotional arc [64–66]. Ranked third in 1880, ‘Emma’ dropped at a gradually increasing rate over the next 90 years to a stable set of low ranks in the 1970s—the decade of ‘Jennifer’—bottoming out at $r=463$ in 1976. After first starting to revive in 1983, ‘Emma’ rapidly rose back to 4th in 2002 and stayed in the top 3 from 2003–2021, six times atop with $r=1$.

Comment 6 by Reviewer 2:

The figures function as standalone illustrations of the concepts and enhance the intuition underlying the measure. I especially like Figure 2, which manages to compress an impressive amount of information into a clear visualization.

Furthermore, the Supplemental Flipbooks are fantastic and I hope the publication of this article is accompanied by a Twitter campaign sharing these as gifs. I'm already excited to explore the allotaxonomy of my own research questions!

Thank you!

As we noted in the paper, we have an online site for the papers, data, and code that we have produced (and will produce) around allotaxonomometry. There are various gifs on display! We will make more!

The top of the code page, for example, shows a video sliding through for season points totals for NBA players, year to year: <https://compstorylab.org/allotaxonomometry/code/>

Our second paper on allotaxonomometry (which introduces probability-turbulence divergence) is also available there. While conceived as an analog of RTD, we show that PTD has a range of very interesting properties of its own. We also show that PTD unifies and generalizes a wide array of existing probability-based divergences.

The main site is here: <https://storylab.w3.uvm.edu/allotaxonomometry/>

Response to Reviewer #3:

Comment 1 by Reviewer 3:

Reviewer #3: This is overall an excellent piece of research. The authors develop an innovative, solid method to compare rankings in systems, and demonstrate its use for word frequencies in social media. The methods are solid, the writing of very quality and the figures as well. It is without hesitation that I recommend this paper for publication as is.

We greatly, greatly appreciate Reviewer 3's extremely positive assessment. Thank you!